Linux: the first second

Alison Chaiken

alison@she-devel.com March 3, 2017





How Linux starts

- What precisely does "off" mean?
- Why does my bootloader have two binaries?
- ACPI vs DTB
- When does the kernel first spawn a second thread?
- What is an initrd?
- How does PID 1 start?
- Out-of-bounds: systemd;
 SecureBoot; fastboot



https://flic.kr/p/6h96xg

THE FOLLOWING CONTENT MAY BE TOO DISTURBING FOR SOME VIEWERS

- -- you will witness a kernel panic and a boot failure;
- -- a NULL pointer will be dereferenced ... successfully!





Applying power





Warm vs. power-on reset

	Clears memory? Restarts clocks?	Pros	Cons	Examples
Power-on Reset	Yes, then reads boot-mode pins.	Won't fail.	Slightly slower.	Plug-in device
Warm Reset	DDR set to 'self-refresh', then reset clocks and jump to stored address.	Faster; retains 'reset reason' and RAM data.	Can fail.	'reboot'; watchdog; JTAG

x86_64: Never genuinely off

Source: Intel

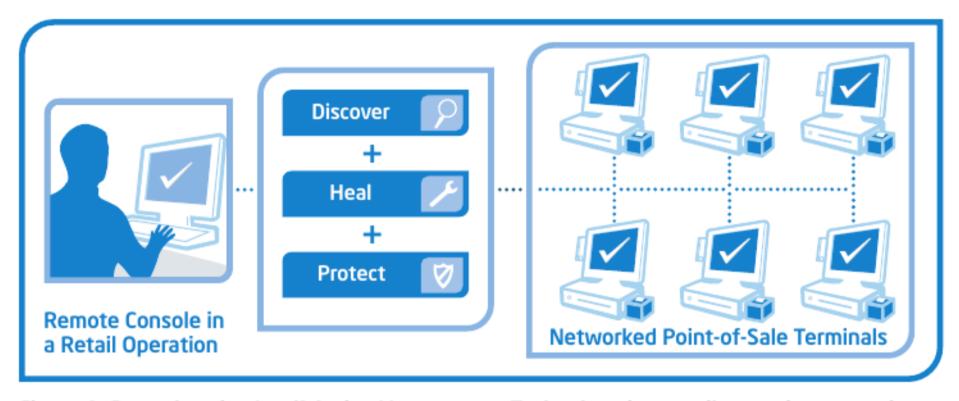
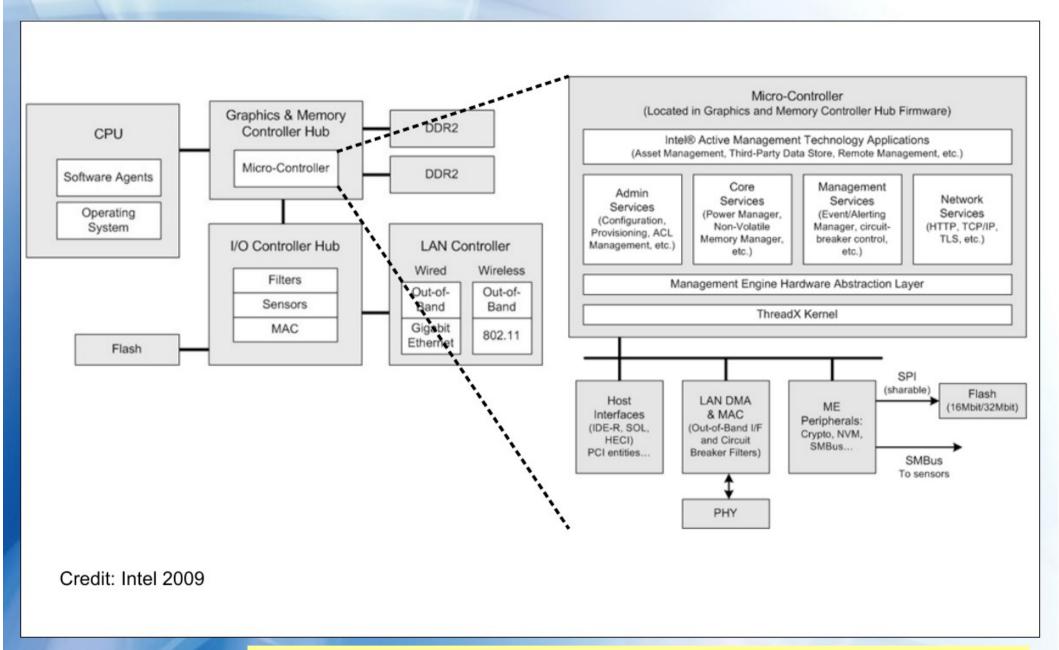


Figure 1. Example using Intel® Active Management Technology in a retail operation to monitor a network of embedded systems even while the enabled systems are powered off.

IPMI: run from Baseboard Management Controller

AMT: run from Graphics and Memory Controller Hub

ME: High-level overview



Source: https://recon.cx/2014/slides/Recon%202014%20Skochinsky.pdf



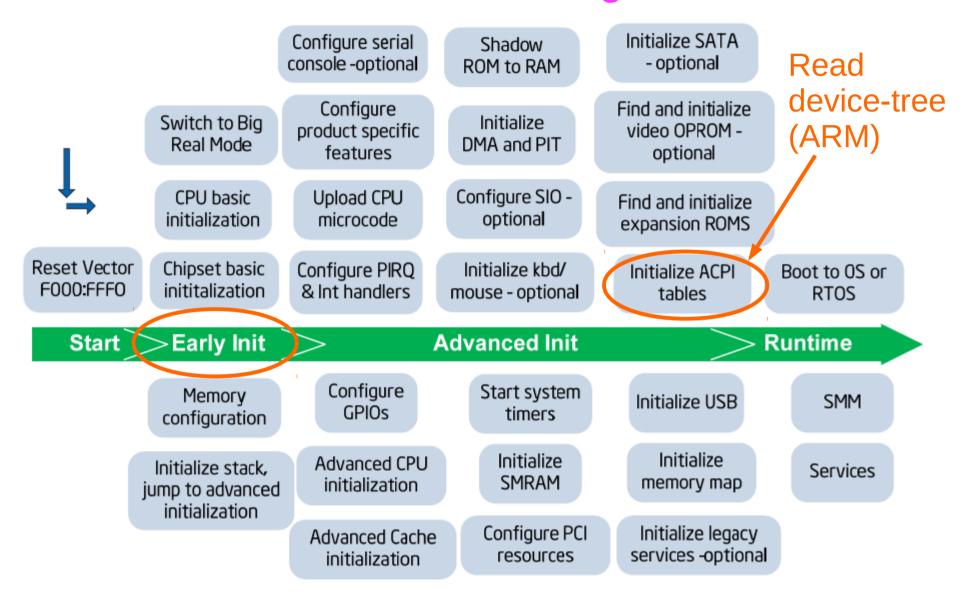


Bootloaders: x86 and u-boot





Bootloaders according to Intel

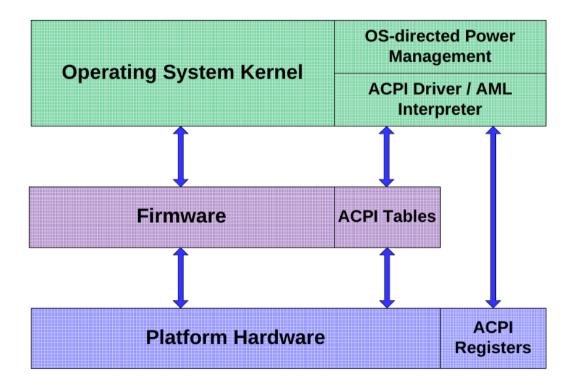


<u>'Shim' bootloader</u> ≈ <u>'Early Init'</u>

- ARM: "SPL", "XLoader" or "MLO" in addition to u-boot.img.
- "Software program loader" separates CPU-specific code.
- Problem: DRAM controller must be initialized.
- Solution: load into SRAM ('OCRAM' in i.MX6, 'l2ram' for TI).
 - Why this works: SRAM (and pNOR) are mapped memory.
- Problem: SRAM is little! (256K on i.MX6, 2 MB on DRA7x).
- Solution: start with a tiny SPL.

Advanced Configuration and Power Interface

Source: Intel



Entertainment: 'sudo acpidump | grep Windows'

Getting more detailed u-boot messages

```
U-boot config:
# Boot timing
CONFIG_BOOTSTAGE=y
CONFIG_BOOTSTAGE_REPORT=y
CONFIG_BOOTSTAGE_FDT=y
```

/* Exiting U-Boot, entering OS */-

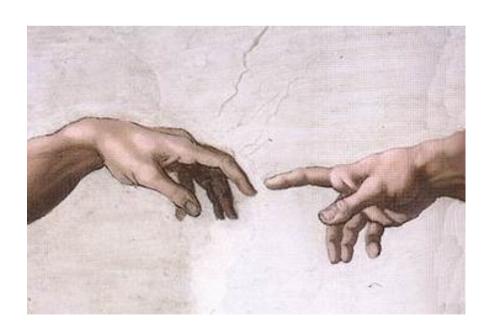
```
Starting kernel ...
Timer summary in microseconds:
              Elapsed
      Mark
                      Stage
                      reset
    71,000 71,000 id=64
    94,000 23,000 id=65
             2,000 main loop
    96,000
 6,137,000 6,041,000 usb start
 22,826,000 16,689,000 id=80
 22,826,000
                   0 eth start
22,826,000
                   0 tftp start
             250,000 ld=82
 23,076,000
                   0 tftp done
 23,076,000
                   0 id=81
 23,076,000
                   0 id=84
 23,076,000
 29,384,000
            6,308,000 id=1
29,384,000
                   O bootm start ***
               18,000 id=3
 29,402,000
29,402,000
                   0 id=2
              186,000 id=5
29,588,000
                   0 id=6
29,588,000
29,588,000
                   0 id=4
              157,000 id=7
 29,745,000
 29,775,000
               30,000 id=15
               28,000 start kernel
 29,803,000
```

<u>Interpretation of bootstages</u>

```
enum bootstage id {
BOOTSTAGE ID START = 0,
BOOTSTAGE ID CHECK_MAGIC, /* Checking image magic */
BOOTSTAGE ID CHECK HEADER, /* Checking image header */
BOOTSTAGE ID CHECK CHECKSUM, /* Checking image checksum */
BOOTSTAGE ID CHECK ARCH, /* Checking architecture */
BOOTSTAGE_ID_CHECK_IMAGETYPE = 5,/* Checking image type */
BOOTSTAGE ID DECOMP IMAGE, /* Decompressing image */
BOOTSTAGE ID KERNEL LOADED, /* Kernel has been loaded */
BOOTSTAGE_ID_DECOMP_UNIMPL = 7, /* Odd decompression algorithm */
BOOTSTAGE ID RUN OS = 15, /* Exiting U-Boot, entering OS */
/* Boot stages related to loading a kernel from an network device */
BOOTSTAGE_ID_NET_CHECKSUM = 60,
BOOTSTAGE ID NET ETH START = 64,
BOOTSTAGE ID NET ETH INIT,
BOOTSTAGE ID NET START = 80,
BOOTSTAGE ID NET NETLOOP OK,
};
```

Passing info from Kernel to Bootloader

- U-boot can pass info in registers or in kernel cmdline.
- Most arches pass "Reset cause" to bootloader.
- mtdoops and ramoops save backtraces in persistent stores specified by device-tree.



mtdoops in ARMv7 Device-tree

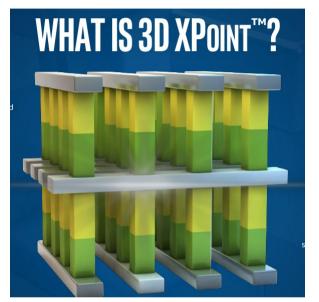
```
flash: m25p80@0 {
        compatible = "sst,sst25vf016b", "jedec,spi-nor";
        spi-max-frequency = <0x1312d00>;
        req = <0x0>;
        \#address-cells = <0x1>:
        \#size-cells = <0x1>:
        mtd@000000000 {
                label = "u-boot.img";
                reg = <0x0 0xc0000>;
        };
        mtd@000C0000 {
                label = "u-boot.env";
                req = <0xc0000 0x2000>;
        };
        mtd@000C2000 {
                label = "splash";
                req = <0xc2000 0x4000>;
        };
        mtd@000C6000 {
                label = "mtdoops";
                reg = <0xc6000 0x13a000>;
        };
```

The coming revolution in non-volatile storage

Source: Micron

Specs:

ArsTechnica



AKA,
'Optane'
by Intel

- 1 Gb non-volatile memory → suspend even for brief inactivity.
- POR will become a rare event.
- Profound security implications.
- Linux drivers: Matthew Wilcox, XIP →DAX

nom notoboon morprotod by companion omarquene app

 Intel Optane Performance: Non-volatile storage medium in the PCIe M.2 format for significant improvements in endurance, performance, and power consumption



ThinkPad X270 (Image credit: Lenovo)





Starting up the kernel





The kernel is an ELF binary

- vmlinux is a regular ELF binary.
 - readelf -e vmlinux
- Extract vmlinux from vmlinuz:
 - <path-to-kernel-source>/scripts/extract-vmlinux \

/boot/vmlinuz-\$(uname -r) > vmlinux



https://flic.kr/p/6xuhiK

?

Quiz:

How do ELF binaries start?

Examining ELF binary start with GDB

(results depend on toolchain)

- Compile your C program with '-ggdb'.
- gdb <some-binary-executable>
- Type 'info files'
- Look for 'Entry point'.

- x86_64:
 - Type 'b *(hex address)'
 - Type 'run'
 - Type 'info functions'

- ARM:
 - Type 'l *(hex address)'
 - Type 'l 1,80'
 - Type 'info functions' or 'info sources'

```
Reading symbols from drm info arm...done.
(adb) info files
Symbols from "/home/alison/gitsrc/drm-tests-0.1/drm info arm".
Local exec file:
        `/home/alison/gitsrc/drm-tests-0.1/drm info arm', file type elf32-littl
        Entry point: 0x893c
        0x00008134 - 0x0000814d is .interp
        0x00008150 - 0x00008170 is .note.ABI-tag
        0x00008170 - 0x00008194 is .note.gnu.build-id
        0x00008194 - 0x00008284 is .gnu.hash
        0x00008284 - 0x000084c4 is .dynsym
        0x0000084c4 - 0x000086ea is .dynstr
        0x000086ea - 0x00008732 is .gnu.version
        0x00008734 - 0x00008754 is .qnu.version r
        0x00008754 - 0x00008764 is .rel.dyn
        0x00008764 - 0x00008814 is .rel.plt
        0x00008814 - 0x00008820 is .init
        0x00008820 - 0x0000893c is .plt
        0x0000893c - 0x00009660 is .text
        0x00009660 - 0x00009668 is .fini
        0x00009668 - 0x00009b48 is .rodata
        0x00009b48 - 0x00009b50 is .ARM.exidx
        0x00009b50 - 0x00009b54 is .eh frame
        0x00011b54 - 0x00011b58 is .init array
        0x00011b58 - 0x00011b5c is .fini array
        0 \times 000011 \text{b5c} - 0 \times 000011 \text{b60} \text{ is .jcr}
        0 \times 000011b60 - 0 \times 000011c60 is .dynamic
        0 \times 00011c60 - 0 \times 00011cc8 is .qot
        0x00011cc8 - 0x00011f58 is .data
        0x00011f58 - 0x00011f60 is .bss
(adb) l *(0x893c)
0x893c is at ../ports/sysdeps/arm/start.S:79.
74
                 .type start,#function
75
        start:
76
                /* Protect against unhandled exceptions. */
77
                .fnstart
78
                 /* Clear the frame pointer and link register since this is the
 frame. */
```

The kernel as PID 0

- Userspace processes need a stack, a heap, STD* file descriptors and an environment to start.
- An ASM constructor "ctor" from crti.o, crtn.o and crt1.0 provided by libgcc allocates these resources.
- Source is start.S.
- Corresponding kernel resources are provided via inline ASM.

Examining kernel start with GDB

```
1 Type 'file vmlinux'. (If zlmage, extract with linux/scripts/extract-vmlinux).
  Type:
      gdb vmlinux
3 Type:
      info files
4 Find 'Entry point'.
  Type:
     l *(hex address)
6 Type
       l 1,80
```

```
0x80dd5e40 - 0x80e8c8f4 is .bss
                                                         Kernel starts in head.S,
(qdb) l *(0x80008000)
0x80008000 is at arch/arm/kernel/head.S:87.
                                                         not start.S.
warning: Source file is more recent than executable.
82
83
        THUMB( badr
                                                @ Kernel is always entered in ARM.
                        r9, 1f
84
        THUMB( bx
                                                @ If this is a Thumb-2 kernel,
                        r9
85
                                                @ switch to Thumb now.
        THUMB( .thumb
86
        THUMB(1:
87
88
       #ifdef CONFIG ARM VIRT EXT
89
                    hyp stub install
               bl
```

What's in head.S?

- Type 'file vmlinux.o'
- Try 'arm-linux-gnueabihf-gdb vmlinux.o'
- Type 'info files'
- Type 'l *(0x0)' <---- actually works!

```
(qdb) l *(0x0),*(0x60)
0x0 is at arch/arm/kernel/head.S:367.
367
                   hyp stub install secondary
       #endif
368
369
               safe svcmode maskall r9
370
371
                                                      @ get processor id
                       p15, 0, r9, c0, c0
               mrc
372
               bl
                      lookup processor type
373
                      r10, r5
                                                       @ invalid processor?
               movs
374
                      r0, #'p'
                                                       @ yes, error 'p'
               moveq
                                   (1) force fixup-able long branch encoding
375
        THUMB( it
                       eq )
376
               bea
                      __error_p
377
378
379
                * Use the page tables supplied from cpu up.
380
                */
                       r4, secondary data
381
               adr
382
               ldmia r4, {r5, r7, r12}
                                                      @ address to jump to after
383
                       lr, r4, r5
                                                       @ mmu has been enabled
               sub
```

The kernel's main() function: highlights of start_kernel()

```
start kernel() {:
                                  "Activate the first processor."
 -boot cpu init(); ◄
 -page address init();
                                        process the device-tree
 -setup arch(&command line); 
 -page alloc init();
 -pr notice("Kernel command line: );
 -parse args("Booting kernel", command line);
 -mm init(); →
                                         setup page tables
                                         and start virtual memory
 -sched init();
 -init IRQ();
                                            All timestamps before
 -init_timers(); timekeeping_init(); 
                                            are [0.000000]
 -rest_init(); 
                       start
                       userspace
                                                                 27
```





About Initrds





What is an initrd anyway?

- 'init ramdisk' = filesystem that is loaded into memory by the kernel before the rootfs mounts.
- Why?
 - To provide a 'rescue shell' in case rootfs doesn't mount.
 - To provide modules that don't fit in zlmage.
 - To provide a safe environment to run agressive tests.
 - To facilitate software updates on devices with limited storage.

What's in an initrd and why?

- Boot into the rescue shell by providing a broken cmdline in /boot/grub/grub.cfg
 - Type 'ls'
- Or try 'Isinitramfs /boot/\$(uname -r)'
- initrd is a gzipped cpio archive:

```
cp /boot/initrd-$(uname -r) /tmp/initrd.gz
gunzip /tmp/initrd.gz
cpio -t < /tmp/initrd</pre>
```

Exploring initramfs

```
(initramfs) ls
bin
          dev
                    init
                               lib64
                                        root
                                                 sbin
                                                           sus
                                                                   var
conf
          etc
                    lih
                              proc
                                                 scripts
                                        run
                                                           tmp
(initramfs) mount
rootfs on / type rootfs (rw)
sysfs on /sys type sysfs (rw, nosuid, nodev, noexec, relatine)
proc on /proc type proc (rw, nosuid, nodev, noexec, relatime)
udev on /dev type devtmpfs (rw,relatime,size=10240k,nr_inodes=1524441,mode=755)
devpts on /dev/pts type devpts (rw, nosuid, noexec, relatine, gid=5, mode=620, ptmxmod
e=000)
tmpfs on /run type tmpfs (rw, nosuid, relatime, size=2442500k, mode=755)
(initranfs) df -h
                                        Used Available Use% Mounted on
                             Size
Filesystem
                                                         0% /dev
                                                 10.0H
                             18.8M
udev
                                       72.0K
                                                  2.3G
                                                         0% /run
                             2.3G
tmpfs
(initramfs)
```

Booting into Rescue Shell

```
Begin: Running /scripts/local-block ... done.
Begin: Running /scripts/local-block ... done.
Begin: Running /scripts/local-block
Begin: Running /scripts/local-block ... done.
done.
Gave up waiting for root device.
                                    Common problems:

    Boot args (cat /proc/cmdline)

   - Check rootdelay= (did the system wait long enough?)
   - Check root = (did the system wait for the right device?)

    Missing modules (cat /proc/modules; ls /dev)

ALERT! /dev/disk/by-uuid/123-seems-unlikely does not exist. Dropping to a shel
 modprobe: module ehci-orion not found in modules.dep
     31.9687481 uhci hcd: USB Universal Host Controller Interface driver
     31.971962] ohci hcd: USB 1.1 'Open' Host Controller (OHCI) Driver
 r
     31.9756251 hidraw: raw HID events driver (C) Jiri Kosina
     31.976950] usbcore: registered new interface driver usbhid
     31.977035] usbhid: USB HID core driver
 BusyBox v1.22.1 (Debian 1:1.22.0-9+deb8u1) built-in shell (ash)
 Enter 'help' for a list of built-in commands.
 bin/sh: can't access tty; job control turned off
  (initramfs) _
```

/scripts/local-block is the function local_block() inside initramfs /scripts/local

OMG! My life is over! (rescue shell tips)

Inhale on a 4-count, then exhale on a 10-count.

- Oh no! 'help' scrolls pages of unreadable crap!
 Relax your jaw. Make circles with your neck.
- Read 'man busybox'.
- 'help | grep' works in busybox.
- Look in /bin and /sbin. There's fsck!!
- You have sed and vi (but not emacs ;-()
- Type 'reboot -f' when you are bored.





Hibernation and Suspension





{Hibernate, Suspend} == Freeze

- Suspend: save state to memory.
- Hibernate: save state to disk (swap).
- Main code is in kernel/freezer.c.
- Freezing of processes preserves state, protects filesystems.
- Each core must save its state separately.
- To watch, append no_console_suspend to bootargs.

How to suspend from CLI

```
root@nitrogen:~# cat /proc/cmdline
console=ttymxc1,115200 ip=172.17.0.1 rootwait rw root=/dev/nfs nfsroot=172.17.0.5:/tftpboot,v3,tcp no console suspend
root@nitrogen:~# echo +30 > /sys/class/rtc/rtc0/wakealarm
root@nitrogen:~# echo mem > /sys/power/state
PM: Syncing filesystems ... done.
Freezing user space processes ... (elapsed 0.001 seconds) done.
Freezing remaining freezable tasks ... (elapsed 0.001 seconds) done.
PM: suspend of devices complete after 78.751 msecs
PM: suspend devices took 0.080 seconds
PM: late suspend of devices complete after 1.506 msecs
PM: noirq suspend of devices complete after 1.500 msecs
Disabling non-boot CPUs ...
CPU1: shutdown
CPU2: shutdown
CPU3: shutdown
Enabling non-boot CPUs ...
CPU1 is up
CPU2 is up
CPU3 is up
PM: noirq resume of devices complete after 1.269 msecs
PM: early resume of devices complete after 1.128 msecs
PM: resume of devices complete after 67.159 msecs
PM: resume devices took 0.070 seconds
Restarting tasks ... done.
atal: SATA link down (SStatus 0 SControl 300)
fec 2188000.ethernet eth0: Link is Up - 1Gbps/Full - flow control rx/tx
```

no_console_suspend in bootargs gives debug output.

How to Hibernate from CLI

```
root@stretch-qemu:~# ls /sys/class/rtc/rtc0
date device max_user_freq power subsystem uevent
dev hctosys name since_epoch time wakealarm
root@stretch-qemu:~# echo +30 > /sys/class/rtc/rtc0/wakealarm
root@stretch-qemu:~# echo disk > /sys/power/state
```

Only works for devices with swap space.

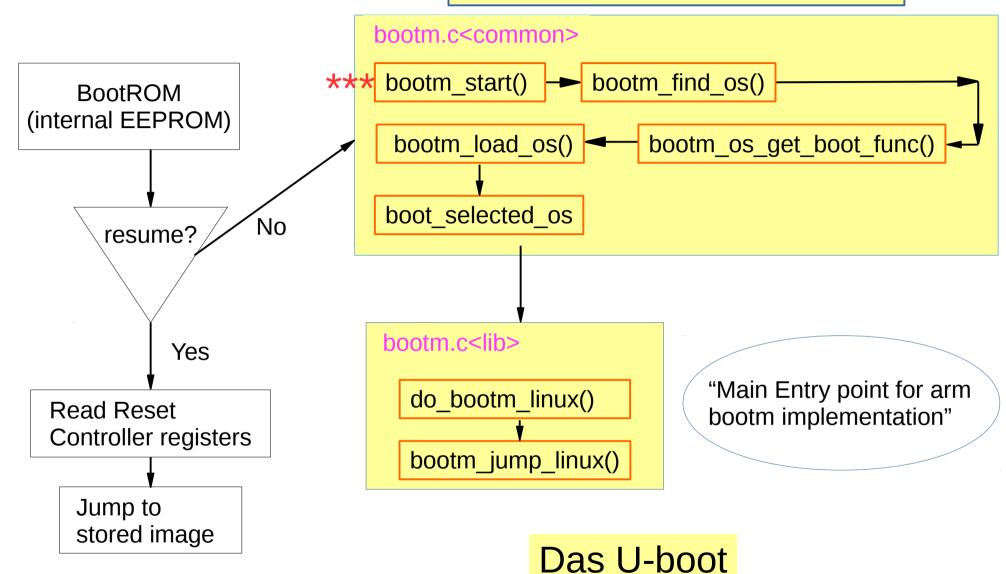
Summary

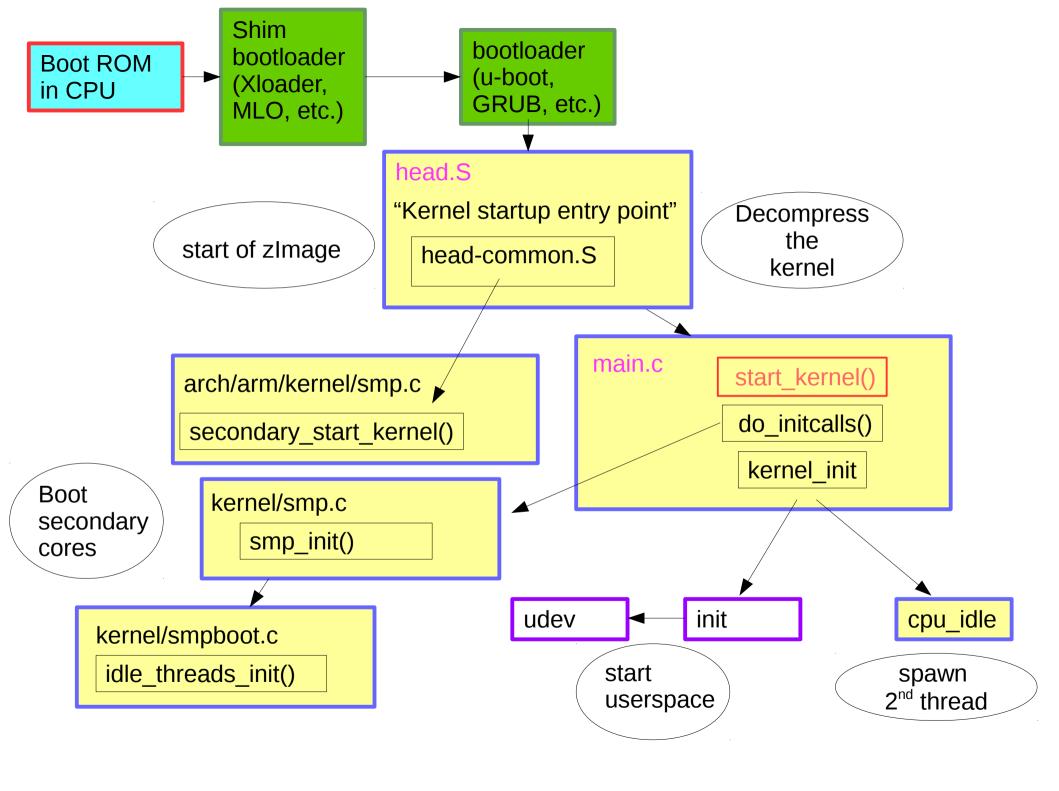
- u-boot's CONFIG_BOOTSTAGES, acpidump and systemd-bootchart provide details about boot.
- Practicing with QEMU is entertaining and low-risk.
- Examine your kernel with GDB and binutils.
- ARM and x86 boot processes are quite distinct.
- High-speed non-volatile memory is about to have a massive impact on Linux.

Major References

- Embedded Linux Primer by Chris Hallinan (book)
- Booting ARM Linux by Russell King and THE LINUX/x86 BOOT PROTOCOL (Documentation/)
- Program startup process in userspace at linux-insides blog
- Matthew Garrett's comprehensive series on UEFI
- Status of Intel Management Engine on various laptops (Coreboot) and servers (FSF)
- All about ACPI talk by Darren Hart, ELCE 2013
- Arch Wiki on hacking ACPI tables

do_bootm_states = u-boot state machine





How to create your own initrd

- Unpack one that already works with gunzip and 'cpio -i'
- Copy in your binary.
- Use gen initramfs.h from kernel source tree:
 - scripts/gen_initramfs_list.sh -o <archive> <path to source>
- Run 'Isinitramfs <archive>' to check the result.
- cp <archive> /boot; edit /boot/grub/grub.cfg
 - CAUTION: your system boots fine, right? You're crazy to mess with the bootloader, you moron.
- Run grub-script-check.

The magnificent result!

```
modprobe: module ehci-orion not found in modules.dep
    32.8051481 uhci_hcd: USB Universal Host Controller
 32.8084021 ohci_hcd: USB 1.1 'Open' Host Controlle
[ 32.812121] hidraw: raw HID events driver (C) Jiri
[ 32.813376] usbcore: registered new interface driver
  32.813459] usbhid: USB HID core driver
 BusyBox v1.22.1 (Debian 1:1.22.0-9+deb8u1) built-in shel
 Enter 'help' for a list of built-in commands.
 /bin/sh: can't access tty; job control turned off
 (initramfs) bin/hello_world.sh
 Never gonna give you up!
  (initramfs)
```

Getting more detailed kernel messages at boot

- Remove 'quiet' from the kernel command line.
- How to keep 'quiet' from coming back:
 - edit /etc/grub.d/10_linux and add:
 export GRUB_DISABLE_SUBMENU=y
 export GRUB CMDLINE LINUX DEFAULT=""

CAUTION: your system boots fine, right? You're crazy to mess with the bootloader, you moron.

Always run 'grub-script-check /boot/grub/grub.cfg' afterwards.

Learning more with systemd-bootchart

- Make sure kernel is compiled with CONFIG_SCHEDSTATS=y.
- 'apt-get install systemd-bootchart'
- Interrupt grub by typing 'e'
- Append 'init=/lib/systemd/systemd-bootchart' to the line that starts with 'linux'
- After boot, open the SVG image in /run/log/ with a browser.

A change in compiling your own kernel

```
kernel/built-in.o
 LD
 CC
          certs/system_keyring.o
make[1]: *** No rule to make target 'debian/certs/benh@debian.org.cert.pem', needed by 'cert
s/x509 certificate list'. Stop.
Makefile:970: recipe for target 'certs' failed
make: *** [certs] Error 2
          rotch_domu linux_stable (version4 & 17)1$ ls a/Pictures

    To: 823107-done@bugs.debian.org

    • Subject: Re: Bug#823107: linux: make deb-pkg fails: No rule to make target 'debian/certs/benh@debian.org.cert.pem'
    • From: Ben Hutchings < ben@decadent.org.uk >

    Date: Sat. 30 Apr 2016 22:50:04 +0200
```

Closing, this is not a bug.

```
You wrote:
[...]
> Should I remove CONFIG SYSTEM TRUSTED KEYS from .config before building
> the kernel? I hope not.
[...]
```

Yes, you must do that. Your custom kernel configuration should be based on the appropriate file provided in linux-source-4.5. These have the CONFIG MODULE SIG ALL, CONFIG MODULE SIG KEY and CONFIG_SYSTEM_TRUSTED_KEYS settings removed so that custom kernels will get modules signed by a one-time key.

Ben.

Appendix: running QEMU

#!/bin/bash ROOTDIR=/home/alison/ISOs HDNAME=debian-testing VERSION=4.9.5

Load kernel via GRUB; console shows in QEMU window. #qemu-system-x86_64 -machine accel=kvm -name \${HDNAME} -boot c -drive file=\$ {ROOTDIR}/\${HDNAME}.raw,format=raw -m 4096 -smp cpus=1 -net nic,model=e1000 -net user,hostfwd=tcp:127.0.0.1:6666-:22 -localtime -serial stdio

Load kernel from external file; console shows in xterm; GRUB doesn't run. qemu-system-x86_64 -machine accel=kvm -name \${HDNAME} -initrd /home/alison/embedded/SCALE2017/kernel/initrd.img-\${VERSION} -kernel /home/alison/embedded/SCALE2017/kernel/vmlinuz-\${VERSION} -boot c -drive file=\$ {ROOTDIR}/\${HDNAME}.raw,format=raw -m 4096 -smp cpus=1 -net nic,model=e1000 -net user,hostfwd=tcp:127.0.0.1:6666-:22 -localtime -serial stdio -append "console=ttyAMA0 console=ttyS0 root=UUID=8e6a1c7e-b3c4-4a37-8e21-56a137c9dded ro"

```
[alison@hildesheim u-boot-imx6 (boundary-v2016.03)]$ file u-boot u-boot: ELF 32-bit LSB shared object, ARM, EABI5 version 1 (SYSV), dynamically linked, interpreter /usr/lib/ld.so.1, not stripped [alison@hildesheim u-boot-imx6 (boundary-v2016.03)]$ arm-linux-gnueabihf-gdb u-boot
```

```
(gdb) info files
Symbols from "/home/alison/gitsrc/u-boot-imx6/u-boot".
Local exec file:
        `/home/alison/gitsrc/u-boot-imx6/u-boot', file type elf32-littlearm.
        Entry point: 0x17800000
        0x178000000 - 0x17852864 is .text
        0x17852868 - 0x1786646e is .rodata
        0x17866470 - 0x1786649c is .hash
        0x178664a0 - 0x1786b25c is .data
        0x1786b25c - 0x1786b268 is .got.plt
        0x1786b268 - 0x1786bdd0 is .u boot list
        0x17877a30 - 0x17877a90 is .dynsym
        0x1786bdd0 - 0x17877a30 is .rel.dyn
        0x1786bdd0 - 0x178b7fd8 is .bss
        0x17877a90 - 0x17877aba is .dynstr
        0x17877abc - 0x17877b3c is .dynamic
        0x17877b3c - 0x17877b4d is .interp
(adb) l *(0x17800000)
0x17800000 is at arch/arm/lib/vectors.S:54.
49
50
        #ifdef CONFIG SYS DV NOR BOOT CFG
51
                        CONFIG SYS DV NOR BOOT CFG
52
        #endi f
53
54
                b
                        reset
55
                ldr
                        pc, undefined instruction
56
                        pc, _software interrupt
                ldr
                        pc, _prefetch abort
57
                ldr
                        pc, data abort
58
                ldr
```

The ARM bootloader

- Read fundamental configuration from fuses, switches and GPIOs.
- Then, for ARM:
 - 1. Setup and initialise the RAM.
 - 2. Initialise one serial port.
 - 3. Detect the machine type.
 - 4. Setup the kernel tagged list. device-tree
 - 5. Load initramfs.
 - 6. Call the kernel image.

Code in the SPL: board_init_f() and jump_to_image_linux()

Where do messages originate?

[54.590327] Starting kernel ...
[54.593459]
Uncompressing Linux... done, booting the kernel.

Linux version 3.0.35-2508-g54750ff (gcc version 4.6.3 #1 SMP PREEMPT

CPU: ARMv7 Processor [412fc09a] revision 10 (ARMv7), cr=10c53c7d

CPU: VIPT nonaliasing data cache, VIPT aliasing instruction cache

Machine: Ereescale i MX 6Quad/Duall ite/Solo Sabre-SD Board

Machine: Freescale i.MX 6Quad/DualLite/Solo Sabre-SD Board Memory policy: ECC disabled, Data cache writealloc CPU identified as i.MX6Q, silicon rev 1.1

PERCPU: Embedded 7 pages/cpu @8c008000 s5440 r8192 d15040 u32768

Built 1 zonelists in Zone order, mobility grouping on. Total pages: 227328

Kernel command line: console=ttymxc1,115200 ip=dhcp rootwait root=/dev/nfs nfsroot=172.17.0.1:/tftpboot/alison/mx6q/fsl-mx6,v3,tcp passed from u-boot

from

kernel

Image, zlmage, ulmage, vmlinux, vmlinuz?

- Image is the raw executable.
- *zImage* is compressed version of Image with prepended uncompression instructions in ASM.
- ulmage is a zlmage with a u-boot header.
- vmlinux is ELF executable containing Image in .text section.
- vmlinuz is a stripped version of vmlinux.